

AI Tournament Judge: design, state and remaining work

How the framework judges student software with independent LLM judge personas and deterministic tooling, what it has done in three events, and exactly what is left before the next version. Written so the owner can return after a break and resume without re-reading the repository.

Doc version 1.0 **Date** 2026-09-24 **Framework** 0.4.0-beta (VERSION)

Repo commit 95ae008 (main) **Owner** Gregg Reed **Companion** docs/design/DESIGN.md

Contents

1. [Summary](#)
2. [What it is and why](#)
3. [Principles](#)
4. [Architecture](#)
5. [Event lifecycle](#)
6. [Per-team judging](#)
7. [Bracket and matchups](#)
8. [Audit, repair and staleness](#)
9. [Privacy and publication](#)
10. [Agents, skills and commands](#)
11. [Data model](#)
12. [The atj CLI](#)
13. [Hooks, tests and release checks](#)
14. [Durable design decisions](#)
15. [History and evidence](#)
16. [Lessons learned](#)
17. [Done and left](#)
18. [Risks](#)
19. [Resume here](#)
20. [Version history](#)

1. Summary

The AI Tournament Judge (ATJ) is a repository-local framework that evaluates student software submissions. Four LLM judge personas score each submission independently against one shared rubric. A consolidator merges their reports. A seeded bracket pairs teams, and each matchup is judged twice in reversed presentation order. The output is a private judging record, a team-facing dossier per team, and approved public summaries. Every number, validation, random draw and state change is produced by a deterministic Python CLI, `atj`. LLM judgment is never described as reproducible.

Question	Answer at commit 95ae008	Source
Released version	0.4.0-beta , 2026-09-18. An Unreleased section holds doc fixes only.	VERSION , CHANGELOG.md

Question	Answer at commit 95ae008	Source
Events completed	Three: <code>sample-mock-2026</code> (scripted fixture), <code>live-trial-2026</code> (first real trial), <code>trial-2-2026</code> (second real trial, completed 2026-09-24)	<code>events/*/status.md</code>
Next version	<code>0.5.0-beta</code> , status <code>proposed</code> . A second trial answering stated hypotheses. The trial ran. The calibration, the findings write-up and the hypothesis answers did not.	<code>docs/0.5.0-beta-plan.md</code>
Open work	30 work items (W1 to W30), 5 carried trial-2 items, 3 plan decisions and one disclosure decision. See section 17 .	<code>docs/0.5.0-beta-plan.md</code> , <code>events/trial-2-2026/audits/final.md</code>
Fit to decide a real award?	No. The framework says so itself: no calibration has ever been produced, and the human attestation acts are unrecorded.	<code>README.md</code> , <code>docs/final-audit.md</code>

2. What it is and why

An LLM can read a codebase and form a view on its quality. On its own that view is not a result anyone can defend: it is not reproducible, it can be steered by text in the submission, and a single model has a single set of blind spots. ATJ wraps the judgment in machinery that makes a disputed result answerable (`README.md`, "What is deterministic").

- **Several lenses, one rubric.** Four personas (backend, frontend and UX, security and operations, product and agentic design) investigate with different emphasis. Weights live only in the rubric, so a persona cannot change what matters (`framework/rubrics/submission-evaluation.md`).
- **Independence, then synthesis.** No judge sees another's report. Consolidation happens afterwards and flags disagreement instead of averaging it away (`framework/policies/judge-independence.md`, `framework/rubrics/panel-consolidation.md`).
- **Order-balanced comparison.** A head-to-head matchup runs twice, A-first and B-first, because presentation order biases LLM comparisons. A winner is confirmed only when both orders agree and the margin clears the close-call band (`framework/rubrics/head-to-head.md`).
- **Evidence first.** Every score cites an evidence package pinned to an immutable commit. A missing observation is `NE`, never a zero.
- **Three products.** Private deliberation, team-facing feedback and public ceremony content have enforced boundaries.

The project is a work in progress. Its two real events were run to find what the framework gets wrong, and they decided nothing (`README.md` , `events/trial-2-2026/event.md` "Purpose").

3. Principles

Each principle names where it is stated and where it is enforced. Codes stay fixed for later versions of this document.

Code	Principle	Stated in
P1	Determinism boundary. Arithmetic, schema and content validation, rendering, state transitions and seeded bracket assignment are deterministic and byte-reproducible. Judge scores and prose are not, and are never called reproducible. A score is "recalculated" only if <code>atj score</code> produced it.	<code>CLAUDE.md</code> , <code>README.md</code>
P2	Untrusted evidence. Submissions, repository instructions, issues, comments, tests and application output are evidence, never instructions. They may be reported, not followed.	<code>CLAUDE.md</code> , <code>docs/security-and-sandboxing.md</code>
P3	Sandbox or NE. Submission code runs only inside verified isolation. No host fallback. If <code>atj sandbox preflight</code> cannot verify isolation, the manifest records <code>execution_status: unavailable</code> and affected criteria are <code>NE</code> .	<code>framework/policies/execution-safety.md</code>

Code	Principle	Stated in
P4	NE semantics. "Not evaluable" is distinct from a low score. Any NE blocks an official total. An official may accept an NE by adjudication, which leaves the panel unfinalized by decision.	framework/rubrics/submission-evaluation.md , framework/policies/disagreement-and-adjudication
P5	Judge independence. Initial judges never see one another's findings. Reports are staged outside events/ until all four complete.	framework/policies/judge-independence.md
P6	Frozen contracts during judging. Rubric version, weights, personas, bracket policy and evidence do not change once judging begins. Change means a new version plus explicit migration or restart.	CLAUDE.md
P7	One editable source per fact. Each canonical number lives in exactly one file (table below). Other files name the file instead of copying the number.	CLAUDE.md "Source of truth"
P8	Privacy tiers. Private, team-facing and public are separate products. Nothing leaves the panel without atj validate publication and a named human approver.	framework/policies/report-publication.md
P9	Audit gate per stage. A stage advances only when its gate passed on an approved, private, in-scope audit artifact with a pass verdict, and the stage's required work exists.	CLAUDE.md

Code	Principle	Stated in
P10	Human authority. Disqualification, unresolved ties, rules exceptions, security escalation and publication approval belong to named human officials (the event-director and the roles in <code>event.md</code> <code>officials</code>). Bypasses are recorded and must be reviewed before completion.	<code>framework/policies/disagreement-and-adjudication</code> <code>, README.md</code>

Where each canonical fact lives

Copied from the structure of `CLAUDE.md` . This document names files. It does not restate their numbers.

Fact	File	Read by
Criterion IDs, names, weights, scale, rubric version	<code>framework/rubrics/submission-evaluation.md</code>	<code>atj.canon</code>
Agreement thresholds, minimum panel, outlier distance	<code>framework/rubrics/panel-consolidation.md</code> front matter	<code>atj.scoring</code>
Comparison values, close-call band, tie-break order	<code>framework/rubrics/head-to-head.md</code> front matter	<code>atj.matchup</code>
Bye policies, default, team range, seeding mode	<code>framework/rubrics/bracket-assignment.md</code> front matter	<code>atj.bracket</code>
Agent and skill versions and digests, superseded versions	<code>framework/personas.md</code>	<code>atj.versions</code>
Artifact shape	<code>framework/templates/</code> , <code>schemas/</code>	<code>atj.reports</code> , <code>atj.schema</code>
Event state	<code>events/<event>/status.md</code>	<code>atj.event</code>

Fact	File	Read by
Framework version	VERSION	atj.__init__

Retired rubric versions are kept in `framework/rubrics/archive/` and loaded by `canon.load_reference()`, so a completed event is recomputed under the contract it was judged under (`CHANGELOG.md` 0.3.0-beta, supersession).

4. Architecture

Markdown with YAML front matter is the reviewable source of truth. JSON Schema and Python provide deterministic validation. Claude Code agents and skills do the judgment and orchestration. Source checkouts live under the git-ignored `workspaces/` (`docs/architecture.md`).




Figure A. Component architecture. Submissions reach only intake and the sandbox. Official numbers come only from `framework/` through `atj.canon`. Public output is rendered only from `public/`.

Repository layout

Path	Contents
atj/	Deterministic tooling, one CLI entry point (<code>atj/cli.py</code> , <code>python3 -m atj</code>)

Path	Contents
framework/rubrics/	submission-evaluation.md , panel-consolidation.md , head-to-head.md , bracket-assignment.md , README.md , archive/ (retired versions)
framework/policies/	evidence-and-citation.md , judge-independence.md , disagreement-and-adjudication.md , execution-safety.md , report-publication.md
framework/templates/	18 artifact templates. Required sections come from each template's ## headings (atj/reports.py)
framework/personas.md	Component registry: version, role, writes contract and content digest per agent and skill, plus an append-only superseded-versions table
schemas/	19 JSON Schema files (18 artifact schemas plus common.schema.json)
events/	_template/ (source for new events), sample-mock-2026/ , live-trial-2026/ , trial-2-2026/
.claude/agents/	7 agents: 4 judges, panel-consolidator , matchup-judge , judging-auditor
.claude/skills/	8 skills (section 10)
.claude/commands/	8 slash commands, each a thin entry into a skill
.claude/hooks/	pre-write.sh , post-write.sh , pre-advance.sh , session-start.sh , user-prompt.sh , lib.sh
tests/	21 test modules, tests/hooks/ , tests/fixtures/ (including bracket-20-team/ and prompt-injection/)
scripts/	Legacy entry points kept as shims over the CLI (scripts/_shim.py); retirement was planned in docs/implementation-detail.md D1
tools/	check_placeholders.py (CI placeholder scan), stage_package_data.py (stages framework data into atj/data/ for wheels)
build_backend.py , pyproject.toml , MANIFEST.in	In-tree build backend that stages data before any wheel, sdist or editable build (CHANGELOG 0.3.0-beta, D31)
workspaces/	Submission checkouts and sandbox state. Git-ignored

Path	Contents
docs/	Operator, security, scoring, tournament, recovery and troubleshooting guides; audit and plan records; this design document under docs/design/

atj modules


In dependency order, from docs/implementation-detail.md section 2, with intake and egress added since.

Module	Responsibility
errors	Stable error codes so hooks, CI and audits match on a code, not prose
frontmatter	Bounded YAML front-matter parsing for untrusted content
canon	Reads rubrics and policies. The only place official numbers enter the system. Fatal if a table total disagrees with its declared total
ids	Derived identifiers and digests, for example ev:<event>:<team>:<commit12>:<digest8>
schema	Offline JSON Schema registry
versions	Framework commit, component registry, version compatibility (require_versions is fatal on mismatch), supersession
scoring	Individual scores, panel consolidation, agreement classes (aligned , material-disagreement , severe-disagreement), outliers, adjudication application, panel integrity
matchup	Order-balanced head-to-head resolution
bracket	Constrained-random seeded assignment with a per-constraint audit, verify that re-derives constraints
event	Stages, gates, units, digests, staleness, overrides, next safe action
publication	Private / team / public boundary, secret and score scan, approval
reports	Artifact validation composed from the above: schema, required sections, placeholders, citations, independence, consolidation re-check
render	Generates score tables, panel blocks and bracket tables into Markdown so numbers are never transcribed by hand
ceremony	Static HTML from approved public artifacts only; refuses a non-conforming artifact

Module	Responsibility
sandbox	Container invocation (no network, read-only mount, capabilities dropped, non-root, resource limits) or an honest refusal
egress	The allowlist proxy: internal network with no route out and no resolver, tinyproxy deny-by-default (docs/security-and-sandboxing.md)
intake	Materialize, pin and enroll one submission; archives get a reproducible snapshot
demo , demo_writer	Generate and check the committed synthetic sample event
cli	The operator surface, plus release-check

5. Event lifecycle

The stage list and gate names below are read from `STAGES` and `STAGE_GATES` in `atj/event.py`. A stage moves forward by exactly one step (`legal_transitions`). Moving backwards is an explicit invalidation through `set_stage`, which marks every unit at or after the target stage stale.



can_advance(event) passes only when

1. gate is "passed", set by atj event gate --audit <artifact> (approved, private, in scope, PASS or PASS WITH ADVIS)
2. STAGE_REQUIREMENTS work exists (manifests, one judgment per expected judge, summaries, bracket.json, r)
3. no unit of this stage incomplete or with audit_result FAIL
4. no unit stale against derive_digests
5. leaving final-audit: every --force-reason bypass reviewed (atj event overrides --review)

Invalidation (backwards)

set_stage to an earlier stage marks every unit at or after it stale. A bypass with --force-reason is allowed while running and recorded in status.overrides; it blocks only the final move to complete until reviewed.

Figure B. The ten stages and nine gates from `atj/event.py`. Gate labels sit on the transition they guard.

Stage	Gate to leave it	Work required (STAGE_REQUIREMENTS)	Typical skill
configuration	configuration-audited	Valid event.md , teams.md , status.md , bracket.md (REQUIRED_FILES)	run-judging-ev /event-init)
intake	roster-frozen	Every team pinned to an immutable commit (checked from evidence onward)	prepare-submis
evidence	evidence-validated	evidence/<team>/manifest.md per eligible team	prepare-submis
initial-judging	judgments-audited	judgments/<team>/<judge>.md for every judge in expected_judges	judge-submissi
consolidation	consolidation-audited	summaries/<team>.md per team	consolidate-ju
bracket	bracket-audited	bracket.json	build-bracket
tournament	tournament-audited	At least one report in matchups/	judge-matchup
dossiers	dossiers-approved	dossiers/<team>.md per team	build-team-dos
final-audit	final-audit-passed	Final audit artifact; all overrides reviewed	audit-judging-
complete	none	Frozen record	

atj event status calls next_action , which checks in order: recorded outputs missing, units stale against current digests, units marked stale, failed audits, pending units in the current stage, then whether the stage can advance. It names the narrowest next step. Completed audited units are never redone (docs/recovery-and-resume.md).

6. Per-team judging




Figure C. One team, intake to consolidation audit. Purple boxes are LLM agents, blue are deterministic `atj` steps, amber are human or safety decisions.

Mechanics that matter

- **Evidence package identity.** `evidence_package_id` is derived from the event, team, commit and a digest of the manifest content. Re-preparing identical evidence yields the same ID. Any change yields a new one, which is what makes staleness detectable (`docs/implementation-detail.md` D2).
- **Bidirectional citations.** The manifest's observation table declares which requirements each observation supports, and `atj validate reports` checks that the requirements table agrees (CHANGELOG 0.3.0-beta). This is a shape check. It cannot tell that a citation points at the wrong content (memory: *validate reports cannot verify citations*).
- **Judge scores.** Each judgment's front matter carries one raw score or `NE` per criterion ID and a confidence. No weight and no total appear in a judgment (`framework/templates/individual-judgment.md`). `atj render judgment` generates the scores table from that front matter.
- **Consolidation.** `atj score` computes the weighted panel result and the agreement class per criterion from the thresholds in `panel-consolidation.md`. The consolidator writes prose. The numbers are generated into the report, and `reports.check_consolidation` reloads judgments, reapplies approved adjudications and fails on any mismatch (`docs/implementation-detail.md` D12).
- **Accepted NE.** `score_override.resolved_score: NE` records that an official reviewed the criterion and the `NE` stands. The panel stays unfinalized by decision. Both trial-2 teams ended

this way.

7. Bracket and matchups

Bracket

`atj bracket build --event-dir <dir> --seed <seed>` places eligible teams by a seeded search with fixed restarts, so the result is a pure function of seed and inputs. The bye policy, default and team range come from `framework/rubrics/bracket-assignment.md`. Each constraint reports `satisfied`, `maximized`, `violated`, `not-applicable` or `infeasible`. `feasible: false` is a legitimate outcome that an official must accept in writing (`docs/implementation-detail.md` D5). `atj bracket verify --event-dir` re-derives every constraint from the roster instead of trusting the file's own audit block, and `--reproduce <roster.json>` redraws from the recorded seed (D6).

Matchup




Figure D. Order-balanced matchup resolution, from `atj/matchup.py resolve`. The band value and the tie-break order are in `framework/rubrics/head-to-head.md`. Criterion-level order disagreement is reported but alone does not block a decisive overall result.

The matchup rubric forbids selecting on initial totals (`head-to-head.md`). Match identifiers come from the drawn bracket, not from a second assignment (D14). A known side effect: `atj bracket advance` writes `winner` into `bracket.json`, which stales the `bracket:draw` unit and makes `--reproduce` fail on every completed bracket (W25, W27).

8. Audit, repair and staleness

The `judging-auditor` agent is the only control that has caught wrong citations, invented counts and false reasons. `atj validate reports` returned zero findings at every stage of `live-trial-2026` while the manifests held nine misdirected citations (CHANGELOG 0.3.0-beta). Every repair round in both real trials introduced new defects, so each repair is re-audited, scoped to its diff (memory: *audits catch their own repairs*).




Figure E. The audit and repair loop, and the unit digests that detect drift. Approving an artifact rewrites its front matter, so an approval after recording stales every unit that digests that file (trial-2 FA7, `events/trial-2-2026/status.md`).

Rules learned from running the loop

- The auditor writes only its report and recommends record, approve, gate and advance commands. It never runs them. It has Bash and once re-recorded a unit itself (memory: *auditor must not run state commands*).
- A repair is derived from the primary source, never copied from the audit's own `repair:` field. Three of five trial-2 bracket rounds introduced a defect that came from the auditor's suggested repair (memory: *audit findings carry their own errors*).
- Record a unit after the last edit to anything it digests. Take ledger times from git. Never estimate a time or force a commit date (trial-2 final audit FB2).
- `atj event gate` reads each finding's `scope` and `blocking`, not the verdict alone (`gate_blocking_findings`, fix-plan D16).

9. Privacy and publication




Figure F. Visibility per directory is `DIRECTORY_VISIBILITY` in `atj/publication.py`. Checks are from `docs/implementation-detail.md` D7.

The repository is public. `visibility: private` is a marker of intent that the git remote does not read. Everything committed under `events/` is world-readable once pushed. For `trial-2-2026` both submissions are public and the event-director approved disclosure (`events/trial-2-2026/overrides/ovr-trial-2-2026-publication-disclosure.md`). For `live-trial-2026`, whose two submissions are private repositories, the disclosure decision is still open (O33).

Known gaps: the gate is a shape check. In trial-2 a short submission commit, a repository URL, an affiliation and an integer margin each passed as CLEAR in a public draft (W22). The template's required `framework_commit` fails the gate at full length (W21). `public_scores` is one event-wide flag with no per-team control (W7).

10. Agents, skills and commands




Figure G. Which skill invokes which agent, from grep of `.claude/skills/*/SKILL.md` and `.claude/commands/*.md`. `/event-init` enters `run-judging-event`. `/event-status` reads `status.md` and reports the next permitted command without running it.

Skill	Invokes agents	Main atj calls	Comm
<code>run-judging-event</code>	none directly	<code>event status/validate/advance</code>	<code>/event</code>
<code>prepare-submission</code>	none	<code>intake, sandbox, validate, event</code>	<code>/team-</code>
<code>judge-submission</code>	four judges, then <code>consolidate-judgments</code> and <code>audit-judging-run</code>	<code>validate</code>	<code>/team-</code>
<code>consolidate-judgments</code>	<code>panel-consolidator</code>	<code>score</code>	via <code>judge-</code>
<code>audit-judging-run</code>	<code>judging-auditor</code>	<code>score, bracket, validate, event, release-check</code>	<code>/event</code>
<code>build-bracket</code>	none	<code>bracket build/verify</code>	<code>/brack</code>
<code>judge-matchup</code>	<code>matchup-judge</code> , once per presentation order	<code>matchup, validate</code>	<code>/match</code>

Skill	Invokes agents	Main atj calls	Comm
build-team-dossier	none (the orchestrator writes)	validate	/team-

Write contract. `framework/personas.md` declares for each component whether it writes its own artifact (a path, and it must hold `Write`) or returns a document for the orchestrator to write (`-` , and it must not hold `Write`). Every component that reads untrusted submission content returns a document. Only `judging-auditor` writes, into `events/<event>/audits/` . `atj release-check` fails in both directions. **Versioning.** Each component carries a version and a content digest of its definition file. An edit fails validation until the version is bumped and `atj personas --refresh` recomputes the digest. Versions are not restated here.

11. Data model

Every artifact is Markdown with YAML front matter. The directory decides the artifact kind and its visibility (`ARTIFACT_KINDS` in `atj/reports.py` , `DIRECTORY_VISIBILITY` in `atj/publication.py`).

Directory	Kind (schema)	Template	Visibility
submissions/	submission-intake	submission-intake.md	private
evidence/	evidence-manifest	evidence-manifest.md	private
runs/	model-run	model-run-record.md	private (JSON records here are not validated, W8)
judgments/	judgment	individual-judgment.md	private
summaries/	consolidated-report	consolidated-team-report.md	private
adjudications/	adjudication	adjudication-report.md	private
matchups/	matchup	matchup-report.md	private
matchup-passes/	matchup-pass	matchup-pass-report.md	private
audits/	audit	audit-report.md	private
calibrations/	calibration	calibration-report.md	private (none ever written)
overrides/	manual-override	manual-override-record.md	private
dossiers/	dossier	team-dossier.md	team

Directory	Kind (schema)	Template	Visibility
public/	public-report	public-matchup-summary.md , event-summary.md	public
event root	event.md (event), teams.md (roster), status.md (status), bracket.md , bracket.json (bracket)	event-configuration.md , team-roster.md , bracket-report.md	no declared visibility, so event.md and bracket.md cannot be approved (W29) and bracket.md is validated by nothing (W15)

Common front matter

Field	Meaning
event_id , team_id , match_id	Identity. Match IDs come from the drawn bracket
commit , evidence_package_id	Pinned submission commit and derived evidence identity
rubric , persona	id@version pins, checked by versions.require_versions . A superseded pin validates as an advisory
framework_commit	Commit of this repository the artifact was produced under
model_requested , model_used , model block	Which model was asked for and which ran, with verified and a note on how identity was established. Scripted fixtures say not-applicable (scripted fixture)
started_at , completed_at	UTC stamps. Nothing validates them as times (configuration F20)
visibility	private , team or public ; must match the directory
approval_state , approved_by , approved_at , approval_note	Set only by atj event approve , which validates first. Templates ship draft
validation_state	unvalidated or valid , written by the approval path
Judgment: judge_id , judge_run_id , scores , confidence	One raw score or NE per criterion. No weights, no total

Field	Meaning
Audit: <code>audit_scope</code> , <code>result</code> , <code>findings</code> (each with <code>severity</code> , <code>scope</code> , <code>blocking</code> , <code>state</code>)	What the gate reads. <code>result</code> must be in <code>PASS_RESULTS</code>

The status ledger (`status.md`)

- **Front matter:** `current_stage`, `blocked`, `stage_gates` (gate to `pending` or `passed`), `gate_evidence` (gate to audit artifact path), `units`, `overrides`, `last_updated`.
- **Unit:** `unit_id` (for example `judging:<team>`), `stage`, `state` (`pending`, `in-progress`, `complete`, `stale`), `input_digest`, `outputs`, `audit_result`, `completed_at`. Re-recording over an unchanged digest keeps the original `completed_at`; a changed digest restamps it; `--completed-at` overrides both (`record_unit`). There is no separate recorded-at time (W28).
- **Body:** stage-gate checkboxes that `atj event validate` compares to the ledger (fix-plan D30), a team progress table, blockers and adjudications, and an append-only activity log of every command, repair and audit with input identity and result.

Adjudications and overrides

- **Adjudication** (`adjudications/`): answers a question the framework returned to a human. Fields include `scope` (the score loader applies only `criterion scope`), `criterion`, `trigger`, `resolution` with `score_override`, `advances_team` for matches, `decided_by`, `decision_authority` (`human-official` may move a total, `agent-substituted` may not, absent means the latter), `substitution_reason`, and append-only `amendments`.
- **Override** (`overrides/`): a recorded human decision that steps over a rule or accepts a deviation, with `scope` and `authorized_by`. Trial-2 wrote overrides for the infeasible affiliation constraint, the publication disclosure, and two model substitutions.
- **Gate bypass:** `atj event advance --force-reason` records an entry in `status.overrides`. `atj event overrides --review <i> --official <role>` records that a person read it, never that it was justified.

12. The atj CLI

Run as `python3 -m atj` or `atj` after `pip install -e .`. Global options: `--root`, `--json`, `--version`. Exit codes: `0` success, `1` validation or audit failure, `2` usage or environment error (`README.md`).

Command	Purpose
<code>rubric</code>	Print the ca rubric it wi use

Command	Purpose
<code>schemas</code>	Self-check shipped schemas
<code>personas [--refresh]</code>	Check the components against agent skill files; refresh digests
<code>event init <id></code>	Create an event from <code>events/_template</code>
<code>event validate <dir></code>	Configure roster, state versions, propose-versions; checkboxes
<code>event status <dir></code>	Stage, gate and the next action
<code>event advance <dir></code>	Advance or refuses when gate or work missing; <code>--force-re</code> records a block
<code>event gate <dir> <gate> <state> --audit <artifact></code>	Record a state result against audit artifacts
<code>event approve <paths> --official <role></code>	Set <code>approval_</code> with validation behind it
<code>event overrides <dir> [--review]</code>	List gate by non-zero event one is unreleased

Command	Purpose
<code>event unit <dir> record stale list</code>	Record or i a ledger ur shows drift
<code>intake <dir> <team> <source> [--ref]</code>	Materialize enroll one submission
<code>sandbox preflight</code>	Report whe verified iso available, a what privi mode
<code>sandbox run <src> <cmd...> [--egress-proxy auto]</code>	Run one cc against a submission isolated
<code>sandbox proxy up down status</code>	Egress pro: makes an € <code>network_a</code> enforceabl
<code>score <judgments-dir json></code>	Consolidat panel unde panel's ow version
<code>render judgment consolidated <paths></code>	Generate s tables into artifacts
<code>matchup <input.json></code>	Resolve an balanced h head
<code>bracket build --event-dir --seed</code>	Draw a reproducib bracket
<code>bracket advance <bracket.json> --match --from</code>	Record a w and carry i

Command	Purpose
<code>bracket verify <bracket.json> --event-dir --reproduce --structure-only</code>	Re-derive constraints from seed, structure o asked by n
<code>validate reports <dir></code>	Validate ev artifact in a
<code>validate publication <artifact dir></code>	Gate before disclosure
<code>ceremony <dir></code>	Render sta ceremony l and dossie approved p artifacts
<code>demo build check</code>	Regenerate verify the committed event
<code>release-check</code>	Every fram level check before a re

13. Hooks, tests and release checks

Hooks (`.claude/hooks/`)

Guard rails, not a security boundary. They run with the operator's permissions and fail open on their own errors (`lib.sh`). Claude Code ignores project `permissions.allow` until the workspace is trusted, so nothing depends on it.

- `pre-write.sh` : blocks event or framework artifact writes on `main` , edits to an active event's frozen contracts after judging begins, and hand-edits of `events/*/public/` .
- `pre-advance.sh` : blocks a stage advance whose gate is pending, and direct host execution of submission code.
- `post-write.sh` : validates what was written; advisory.

- `session-start.sh`, `user-prompt.sh`: report framework health, branch and each event's next safe action.

Tests

21 modules in `tests/`. Regression files name the defect each test locks down (`test_live_trial_regressions.py`, `test_tier1..3_regressions.py`, `test_final_audit_regressions.py`, `test_release_audit_regressions.py`, `test_audit_regressions.py`). Live egress tests need a container runtime and `ATJ_LIVE_SANDBOX=1` and are skipped by default. Bracket tests sweep every size, policy and many seeds. A committed-audit regression in `test_tier1_regressions.py` was adjusted when trial-2 completed with the `findings:` audit shape (commit `381920a`).

atj release-check

At `95ae008` it reports PASS on: rubric, schemas, personas, templates, claude components, single-source, template schemas, version-skew, packaging, version archive, write contracts, signed approvals (completed events frozen), sample event. CI also installs the wheel into a fresh virtualenv and runs `release-check` from `/tmp`. The full validation chain is in `README.md` "Validation" and `docs/release-checklist.md`.

14. Durable design decisions

Coded `DD` here to avoid collision with the `D` codes already used in `docs/implementation-detail.md`, `docs/framework-fix-plan.md` and `docs/0.5.0-beta-plan.md`. Each names its source.

Code	Decision	Source
DD1	The rubric Markdown is the only editable source of official numbers. A second copy anywhere fails the build.	implementation-detail D1
DD2	Identifiers are derived from content, never assigned, so staleness is detectable.	implementation-detail D2
DD3	Version incompatibility is fatal. Retired versions are recorded append-only, so bumps do not falsify completed events.	implementation-detail D3; fix-plan D28
DD4	The close-call band is a floor. Events may widen it, never narrow it, and never from a matchup input file.	implementation-detail D4

Code	Decision	Source
DD5	Bracket placement is a seeded search with a reported status per constraint. Infeasible is an outcome, not a hidden failure.	implementation-detail D5
DD6	Verification re-derives from inputs. It never trusts an audit block inside the file under test.	implementation-detail D6
DD7	The publication gate fails closed and the ceremony reads only <code>public/</code> .	implementation-detail D7
DD8	Execution is isolated or unavailable. No host fallback. A network allowlist requires the egress proxy.	implementation-detail D8; CHANGELOG 0.4.0-beta
DD9	Prompt-injection resistance is claimed only for mechanical properties (scores, rubric, visibility, tools), not for judge prose.	implementation-detail D9
DD10	The sample event's scores are scripted so CI is reproducible without a model, and every artifact says so.	implementation-detail D10
DD11	Hooks are guard rails and fail open.	implementation-detail D11
DD12	Written official numbers are re-derived and compared, not trusted.	implementation-detail D12
DD13	A gate needs an approved, in-scope audit artifact and the stage's work behind it.	implementation-detail D13
DD14	A winner advances only from a confirmed result or an approved adjudication naming it.	implementation-detail D14
DD15	Components that read untrusted content return documents; the orchestrator writes. Declared in the registry and checked both ways.	<code>framework/personas.md</code> ; fix-plan D1
DD16	Evidence stays frozen once judging begins, even to rescue an <code>NE</code> . Check what a missing total blocks before paying to recover it.	<code>CLAUDE.md</code> ; memory <i>frozen evidence beats a better score</i>
DD17	Each trial states its hypotheses in <code>event.md</code> "Purpose" before judging. Planted defects and contamination never enter an official event.	0.5.0-beta-plan D2; <code>events/trial-2-2026/event.md</code>

Code	Decision	Source
DD18	Calibration runs before any award-deciding use.	0.5.0-beta-plan D1; docs/final-audit.md

15. History and evidence

Releases (CHANGELOG.md)

Version	Date	What it was
0.1.0-alpha		Scaffold: rubrics, policies, templates, agent and skill definitions, five scripts, eleven tests. Its documents described controls that did not exist (docs/release-readiness-audit.md).
0.2.0-beta	2026-09-16	First release able to run a complete supervised mock event: atj , canonical data model, schemas, component registry, real sandbox or refusal, publication gate, scripted sample event.
0.3.0-beta	2026-09-18	The release that ran live-trial-2026 and fixed its 31 defects (docs/framework-fix-plan.md). Version supersession, event overrides , render , matchup-pass kind, bidirectional citations, accepted NE , decision_authority , write contracts, build backend.
0.4.0-beta	2026-09-18	Egress proxy verified against a live runtime; per-criterion citation check; completion refused while a gate bypass is unreviewed.
Unreleased		Documentation fact corrections only.

Events

Event	What it exercised	Outcome at public level
sample-mock-2026	Synthetic: four invented teams, scripted scores. Demonstrates material and severe disagreement, an outlier, an initial <code>NE</code> , three adjudications, a close-call matchup with no winner, an order-balanced matchup, a failed then repaired consolidation audit. CI regenerates and checks it (<code>atj demo check</code>).	Complete. Public summaries in <code>events/sample-mock-2026/public/</code> name Lumen as final winner (fixture data).
live-trial-2026	First real trial, the event behind 0.3.0-beta: two real applications, real container isolation, real model judgment, nine audited gates. Found 31 defects, the last two by running the documentation (D29, D31).	Complete 2026-09-18. <code>public/</code> is empty; nothing was published.
trial-2-2026	Second real trial at 0.4.0-beta, framework commit <code>ea0db07</code> at init. Two teams from the operator's own public repositories (<code>beekeeper-lab/ScribeVault</code> , <code>beekeeper-lab/ai-security-demos</code>), 11 isolated runs, no network, no model calls by submissions. Hypotheses in <code>event.md</code> : H6 (submission agent instructions), the NE-versus-low-score boundary on a submission with no tests, and a desktop GUI surface. First adjudications, first overrides, first approved public artifact.	Complete 2026-09-24 (PR #26, main <code>95ae008</code>). <code>public/mu-final-01.md</code> : AI Security Demos advances; both presentation orders agreed. Both panels unfinalized by accepted <code>NE</code> . <code>public_scores: false</code> .

trial-2-2026 audit record

Every stage ended `PASS WITH ADVISORIES` . Round counts are from each audit's front matter `audit_rounds` or `audit_scope` , and the activity log in `events/trial-2-2026/status.md` .

Stage audit	Rounds	First round	Notable	P R
audits/configuration.md	3	FAIL	Repairs introduced new defects each round; the auditor invented timestamps (F19); F20, F23 left for the framework	#18
audits/intake.md	3	FAIL	18 findings, 14 repaired; F7 open by choice	#19
audits/evidence.md	3	FAIL	34 findings over three rounds; F14 (runs/*.json unvalidated) left open	#20
audits/judgments.md	3	FAIL	About 120 citations followed. Two judges reported a secrets-filtered file as absent (W11). H6 killed: no judge followed submission instructions	#21
audits/consolidation.md	4	FAIL	Numbers exact; prose inflated corroboration in one direction; repairs introduced twelve more	#22
audits/bracket.md	5	FAIL	Two same-affiliation teams made the draw infeasible; accepted by override. 45 findings across five rounds	#23
audits/tournament.md	5	PASS WITH ADVISORIES	Round four FAIL (a stale unit after an erratum). W21 to W25 found	#24
audits/dossiers.md	4	PASS WITH ADVISORIES	Fix ac513ed: event approve made every public artifact unapprovable. Summary errata recorded in the ledger (W26)	#25
audits/final.md	6	PASS WITH ADVISORIES	No blocking or major finding. FA2, FA7, CF11 settled by the event-director. W27 to W30 found	#26

16. Lessons learned

Code	Lesson	Evidence
L1	A green <code>atj validate reports</code> means well-formed, not true. The LLM audit is the only citation control.	CHANGELOG 0.3.0-beta; memory <i>validate reports cannot verify citations</i>
L2	Every repair round introduces defects. Re-audit the repair diff, always.	Both trials; memory <i>audits catch their own repairs</i>
L3	An audit finding can be wrong. Derive a repair from the source, never from the audit's repair text.	trial-2 consolidation P1, bracket rounds
L4	One wrong number in an evidence manifest becomes several wrong judge findings, because judges trust the package.	live-trial team-podcast manifest
L5	Consolidation errors run one way: corroboration inflated, single-judge concerns dropped. Count judges per claim.	trial-2 consolidation, final audit advisory
L6	Judges report tool blindness as fact. An absence claim needs <code>git ls-files</code> .	trial-2 judgments F1, W11
L7	A causal claim can be false while every number around it is right. Check stated reasons with a command and a control.	trial-2 configuration, memory <i>repository visibility is checkable</i>
L8	Running the documentation finds defects that reading the code does not.	live-trial D29, D31
L9	Approving an artifact rewrites it and states every unit that digests it. Approve before recording units.	trial-2 final audit FA7 rows
L10	The auditor must not run state commands.	trial-2 tournament round four
L11	Take ledger times from git. Never estimate or force a commit date.	trial-2 final audit FB2

17. Done and left

Done

Area	State	Evidence
Releases	done 0.1.0-alpha to 0.4.0-beta, tags v0.3.0-beta, v0.4.0-beta	CHANGELOG.md, git tag
All ten stages exercised with a real model	done twice	events/live-trial-2026/status.md, events/trial-2-2026/status.md
Live-trial defects	done all 31 fixed	docs/framework-fix-plan.md
Earlier plans	closed no unchecked box in docs/implementation-plan.md	docs/0.5.0-beta-plan.md "What this version is"
Adjudication, override, public artifact, egress-free sandbox runs in a real event	done first time in trial-2	events/trial-2-2026/adjudications/, overrides/, public/
H6 (submission agent instructions cannot steer the panel)	answered concern refuted in one trial	events/trial-2-2026/audits/judgments.md
Release checks and tests	pass at 95ae008	<u>section 20</u>

0.5.0-beta hypotheses

ID	Hypothesis	State
H1	Independence detector misses a paraphrasing judge	not run needs calibration (W1, W2)
H2	A recorded stage audit is a performed one	not run needs calibration (W1, W2)
H3	Bracket and tournament hold at a real size	not answered trial-2 had two teams; the plan says below four H3 is not answered
H4	Publication boundary holds under a real approval	partial one public artifact approved; W21, W22 show the gate is a shape check
H5	Model judgment is stable enough to mean something	not run no calibration report exists
H6	Submission agent instructions cannot steer the panel	answered killed (no steering)

Open items

O-codes are this document's. Each maps to its source ID. "Blocks next" means blocks shipping 0.5.0-beta as defined in "What 0.5.0-beta ships" of docs/0.5.0-beta-plan.md. W8 to W30 are marked "framework scope, no trigger" in the plan, meaning they land only when a change is made for another reason; they block nothing by that rule, but most are defects the next event will hit again.

Code	Source	Item	Owner	Blocks next
O1	W1	Calibration runnable end to end: skill or command, report, validation	framework	yes
O2	W2	Marked control run that cannot be mistaken for official judgment	framework	yes (before H1, H2)
O3	W3	Independence detection beyond wording overlap	framework	conditional on H1
O4	W4	Whatever H3 breaks in bracket, matchup, advancement	framework	conditional on H3

Code	Source	Item	Owner	Blocks next
O5	W5	Whatever H4 breaks in publication	framework	conditional; W21, W22 are its first entries
O6	W6	Amend <code>docs/final-audit.md</code> advisories in place with trial evidence	framework	yes
O7	W7	Per-team publication restriction or a validated refusal to mix	framework	only if a roster team's source is not public
O8	W8 = evidence F14	Validate JSON run records under <code>runs/</code>	framework	no
O9	W9	Artifact must end where its template ends	framework	no
O10	W10	An NE split is a disagreement; outlier and minimum-panel checks per criterion	framework	no
O11	W11	Judges must use <code>git ls-files</code> before asserting absence (persona change)	framework	no
O12	W12	One amendment form for every artifact class	framework	no
O13	W13	Adjudication template yields a record the schema rejects and the loader ignores	framework	no
O14	W14	NE-disposing adjudication must record a rationale	framework	no
O15	W15	<code>bracket.md</code> validated and digested	framework	no
O16	W16	<code>officials</code> schema declares every human-reserved decision	framework	no
O17	W17	Infeasible bracket cannot advance without its override	framework	no

Code	Source	Item	Owner	Blocks next
O18	W18	Bracket report template matches the renderer	framework	no
O19	W19	<code>bracket verify --reproduce</code> reports a bad roster instead of crashing	framework	no
O20	W20	Bracket policy states the affiliation rule the code enforces (frozen-contract change)	Gregg decides, framework implements	no
O21	W21	Publication gate distinguishes framework commit from submission commit	framework	no (H4 evidence)
O22	W22	Gate catches short commits, repo URLs, affiliation, integer margins	framework	feeds H4
O23	W23	Single-source check reads Markdown tables	framework	no
O24	W24	Matchup-result exemption verifies content, not key presence	framework	no
O25	W25	Recording a winner states <code>bracket:draw</code>	framework	no
O26	W26 (CF4)	Errata route for an audited summary	framework	no
O27	W27 (FA1)	Bracket reproduction ignores recorded winners	framework	no
O28	W28 (FA9)	Units carry a separate <code>recorded_at</code>	framework	no
O29	W29	Event-root artifacts (<code>event.md</code> , <code>bracket.md</code>) can be approved	framework	no
O30	W30	<code>event unit record</code> prints "kept" only when it kept	framework	no
O31	intake F7 (CF7)	Roster template has more identity fields than any roster; decide, for example <code>event init</code> writes full front matter for new rosters	Gregg	no

Code	Source	Item	Owner	Blocks next
O32	configuration F10, F20, F23 (CF8 to CF10)	F10 absolute home paths in provenance; F20 nothing validates timestamps as times; F23 <code>atj intake</code> writes <code>model_used: not-applicable</code> unconditionally	Gregg (per <code>audits/final.md</code> carried table)	no
O33	memory <i>public repo holds private judgments</i>	<code>live-trial-2026</code> disclosure: private submissions' judgments are world-readable. Options put: leave, scrub, make repository private	Gregg	blocks any private-source roster
O34	plan Q2, Q3, Q4	No recorded resolution in the plan. Trial-2 used two teams, its own id, and no network. Decide whether trial-2 counts as the 0.5.0-beta event or a further event at four or more teams is needed for H3 and the egress proxy	Gregg	yes
O35	plan "What 0.5.0-beta ships"	Findings document in the shape of <code>docs/framework-fix-plan.md</code> for trial-2, and each hypothesis answered in writing	framework	yes
O36	1.0 line	Turn the eight human-only acts in <code>docs/final-audit.md</code> into recorded attestations. Explicitly out of 0.5.0-beta	Gregg	no (later)

18. Risks

Code	Risk	Source
R1	A trial finds little because submissions are alike. The no-AI and failing-tests cases are still unfilled.	0.5.0-beta-plan R1
R2	A planted defect contaminates the real record. Nothing planted goes near an event until W2 lands.	0.5.0-beta-plan R2
R3	Separation of duties cannot be tested while one person holds every authority. Trial-2 approvals were run by the orchestrator on delegation (CF11, recorded, not provable).	0.5.0-beta-plan R3; <code>audits/final.md</code>
R4	Roster cost: roughly <code>4 x teams + 2 x matches</code> model invocations, and each stage took several audit rounds in trial-2.	0.5.0-beta-plan R4; <code>README.md</code>
R5	A third-party public project did not ask to be judged. Publishing about it is decided at the approval gate.	0.5.0-beta-plan R5
R6	The public repository publishes private judgments (O33).	memory
R7	Independence is procedural and detected, not prevented; a well-paraphrased contamination would pass.	implementation-detail limitations 11, 13
R8	Trial-2 ran some stages on a model other than the one configured, accepted by override. Model identity is recorded truthfully but cross-stage comparability is weaker.	<code>events/trial-2-2026/overrides/</code>

19. Resume here

In order. Each step is small and names its file.

1. **A1.** Branch from `main` and confirm the baseline: `python3 -m pytest tests/ -q && python3 -m atj release-check && python3 -m atj demo check`. Run `python3 -m atj event status events/trial-2-2026` and expect `complete` with nothing stale.

2. **A2.** Read `docs/design/DESIGN.md`, then `docs/0.5.0-beta-plan.md` "Hypotheses", "Work items" and "Open decisions". Do not re-read the audits; section 15 summarises them.
3. **A3.** Decide O34: does trial-2 count as the 0.5.0-beta event, and is a further event at four or more teams needed for H3? Record the answer in the plan's "Decisions taken".
4. **A4.** Decide O33 (live-trial disclosure) before putting any private-source submission on a roster.
5. **A5.** Write O35: the trial-2 findings document and written answers for H4 and H6, and mark H1, H2, H3, H5 as not yet run.
6. **A6.** Land the cheap mechanical fixes as one framework branch before any new event: W25, W27, W28, W30 (unit and bracket bookkeeping), W13, W14 (adjudication template), W29, W15 (event-root artifacts), W19. Each has a named file and line in the plan.
7. **A7.** Build W1 and W2, then run the calibration (H5 first, then H1 and H2 inside it) per plan D1.
8. **A8.** Settle O31 and O32 when touching intake or templates. Amend `docs/final-audit.md` advisories (W6). Then cut `0.5.0-beta` and add a row to this document's version history.

20. Version history

Doc version	Date	Framework	Repo commit	Change
1.0	2026-09-24	<code>0.4.0-beta</code>	<code>95ae008</code>	First versioned design document, written after <code>trial-2-2026</code> completed

Verification at writing: `python3 -m atj release-check` and `python3 -m pytest tests/ -q` were run with this file and `docs/design/DESIGN.md` in the tree on 2026-09-24: release-check PASS on every check; pytest 520 passed, 5 skipped (live egress), 298 subtests passed.

Self-contained file. No network requests. Canonical numbers are deliberately absent; follow the named files. Markdown companion: `docs/design/DESIGN.md`.